

Flight Delay (JAN 2015)

```
In [ ]: import pandas as pd
import numpy as np
import scipy
import scipy.stats as st
import matplotlib.pyplot as plt
# allows for plots to always be displayed in Jupyter notebook
%matplotlib inline
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [ ]: # Upload airlines.csv, airports.csv, flights.csv
from google.colab import files
csvFiles = files.upload()
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving airlines.csv to airlines.csv
Saving airports.csv to airports.csv
Saving flights.csv to flights.csv

```
In [ ]: # Read file
airlines_data = pd.read_csv('airlines.csv')
airports_data = pd.read_csv('airports.csv')
flights_data = pd.read_csv('flights.csv')
```

```
In [ ]: # Visualize Data
airlines_data
```

Out[]:

	IATA_CODE	AIRLINE
0	UA	United Air Lines Inc.
1	AA	American Airlines Inc.
2	US	US Airways Inc.
3	F9	Frontier Airlines Inc.
4	B6	JetBlue Airways
5	OO	Skywest Airlines Inc.
6	AS	Alaska Airlines Inc.
7	NK	Spirit Air Lines
8	WN	Southwest Airlines Co.
9	DL	Delta Air Lines Inc.
10	EV	Atlantic Southeast Airlines
11	HA	Hawaiian Airlines Inc.
12	MQ	American Eagle Airlines Inc.
13	VX	Virgin America

In []:

```
# Visualize Data
airports_data
```

Out[]:

	IATA_CODE	AIRPORT	CITY	STATE	COUNTRY	LATITUDE	LONGITUDE
0	ABE	Lehigh Valley International Airport	Allentown	PA	USA	40.65236	-75.44040
1	ABI	Abilene Regional Airport	Abilene	TX	USA	32.41132	-99.68190
2	ABQ	Albuquerque International Sunport	Albuquerque	NM	USA	35.04022	-106.60919
3	ABR	Aberdeen Regional Airport	Aberdeen	SD	USA	45.44906	-98.42183
4	ABY	Southwest Georgia Regional Airport	Albany	GA	USA	31.53552	-84.19447
...
317	WRG	Wrangell Airport	Wrangell	AK	USA	56.48433	-132.36982
318	WYS	Westerly State Airport	West Yellowstone	MT	USA	44.68840	-111.11764
319	XNA	Northwest Arkansas Regional Airport	Fayetteville/Springdale/Rogers	AR	USA	36.28187	-94.30681
320	YAK	Yakutat Airport	Yakutat	AK	USA	59.50336	-139.66023
321	YUM	Yuma International Airport	Yuma	AZ	USA	32.65658	-114.60597

322 rows × 7 columns

◀

▶

```
In [ ]: # Visualize Data
        flights_data
```

Out[]:

	MONTH	DAY	AIRLINE	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE
0	1	1	AS	ANC	SEA	5
1	1	1	AA	LAX	PBI	10
2	1	1	US	SFO	CLT	20
3	1	1	AA	LAX	MIA	20
4	1	1	AS	SEA	ANC	25
...
469963	1	31	B6	JFK	BQN	2359
469964	1	31	DL	SEA	DTW	2359
469965	1	31	F9	DEN	TPA	2359
469966	1	31	F9	DEN	ATL	2359
469967	1	31	UA	ANC	DEN	2359

469968 rows × 27 columns

In []: *# The column names as indices.*
The two brackets are required because you are passing a list of columns
loc_data = airports_data[['LATITUDE', 'LONGITUDE']]
loc_data


Out[]:

	LATITUDE	LONGITUDE
0	40.65236	-75.44040
1	32.41132	-99.68190
2	35.04022	-106.60919
3	45.44906	-98.42183
4	31.53552	-84.19447
...
317	56.48433	-132.36982
318	44.68840	-111.11764
319	36.28187	-94.30681
320	59.50336	-139.66023
321	32.65658	-114.60597

322 rows × 2 columns

In []: `fig, ax = plt.subplots()`
s=0.1 specifies the size
`ax.scatter(loc_data.LONGITUDE, loc_data.LATITUDE, s=0.1)`

```
ax.set_title('USA Airport Locations')  
ax.set_xlabel('Longitude')  
ax.set_ylabel('Latitude');  
fig.set_dpi(200)
```



```
In [ ]: # [['ARRIVAL_DELAY', 'DEPARTURE_DELAY']] shows dataframe of delays  
df = flights_data[['DESTINATION_AIRPORT', 'ARRIVAL_DELAY', 'DEPARTURE_DELAY']]  
# adding .DESTINATION_AIRPORT at the end creates the rows to be CALLABLE  
airports = df.DESTINATION_AIRPORT
```

```
In [ ]: # The following calculates the mean of all columns  
# and groups all similar names i.e. airports under the same group  
delay_mean = df.groupby(airports).mean()  
  
delay_mean
```

Out[]:

	ARRIVAL_DELAY	DEPARTURE_DELAY
DESTINATION_AIRPORT		
ABE	4.166667	4.044586
ABI	12.361233	10.643172
ABQ	5.538857	8.567949
ABR	-3.716667	7.524590
ABY	6.812500	9.753086
...
VPS	3.793566	7.566489
WRG	9.300000	15.393443
XNA	14.845528	14.817568
YAK	-5.525424	-3.466667
YUM	6.484848	3.866667

312 rows × 2 columns

```
In [ ]: # .sort_values will sort alphabetically the column specified in a csv file
key_array = flights_data.sort_values(['DESTINATION_AIRPORT'])
# adding .IATA_CODE at the end creates the rows to be CALLABLE
key_array = key_array.DESTINATION_AIRPORT

# .values allows for the mean values of arrival delay to be callable
mean_arr_array = delay_mean[['ARRIVAL_DELAY']].values[:,0]
# .around will round up values throughout the entire array
mean_arr_array = np.around(mean_arr_array, decimals=2)

# .values allows for the mean values of departure delay to be callable
mean_dep_array = delay_mean[['DEPARTURE_DELAY']].values[:,0]
# .around will round up values throughout the entire array
mean_dep_array = np.around(mean_dep_array, decimals=2)
```

```
In [ ]: # .sort_values will sort alphabetically the column specified in a csv file
key_array = flights_data.sort_values(['DESTINATION_AIRPORT'])
# adding .IATA_CODE at the end creates the rows to be CALLABLE
key_array = key_array.DESTINATION_AIRPORT
# drops any duplicates
key_array = key_array.drop_duplicates(keep = 'last')
key_array
```

```
Out[ ]: 278297    ABE
        117844    ABI
        356705    ABQ
        9034     ABR
        303884    ABY
        ...
        70454     VPS
        150236    WRG
        148691    XNA
        461870    YAK
        106376    YUM
Name: DESTINATION_AIRPORT, Length: 312, dtype: object
```

```
In [ ]: # airports_data['IATA_CODE'] and key_array should have the same # of airports
        # printing both shows discrepancies of the two lists
        airports_data['IATA_CODE']
```

```
Out[ ]: 0      ABE
        1      ABI
        2      ABQ
        3      ABR
        4      ABY
        ...
        317    WRG
        318    WYS
        319    XNA
        320    YAK
        321    YUM
Name: IATA_CODE, Length: 322, dtype: object
```

```
In [ ]: # air_flightsdata and air_airportsdata creates a list of the airports from the csv file
        air_flightsdata = flights_data['DESTINATION_AIRPORT'].sort_values()
        air_airportsdata = airports_data['IATA_CODE']
        # coordinate data of airports
        longitude_coor = airports_data['LONGITUDE']
        latitude_coor = airports_data['LATITUDE']
```

```
In [ ]: # removal of duplications
        air_flightsdata = air_flightsdata.drop_duplicates()
```


```
In [ ]: # the variables are created into lists
        # a for loop goes through the data in the airport csv file and checks if
        # it is amongst the flights csv file. Any airport that does not match has the index stored
        # .count(example) checks if the example is found in the list interested in
        removeindex = []
        air_airportsdata = list(air_airportsdata)
        air_flightsdata = list(air_flightsdata)
        longitude_coor = list(longitude_coor)
        latitude_coor = list(latitude_coor)
        for i in range(len(air_airportsdata)):
            if air_flightsdata.count(air_airportsdata[i]) < 1:
                removeindex.append(i)
        removeindex
```

```
Out[ ]: [5, 13, 29, 30, 89, 133, 147, 161, 219, 318]
```


```
In [ ]: # the loop uses the stored index from the previous code block
# removing airports that are not shared between the two csv sheets
for i in range(len(removeindex)):
    k = removeindex[i] - i
    air_airportsdata.remove(air_airportsdata[k])
    longitude_coor.remove(longitude_coor[k])
    latitude_coor.remove(latitude_coor[k])
```

```
In [ ]: # dictionary created for airport stats
airport_stats = {}
i = 0
for key in key_array:
    airport_stats[key] = {'Mean Arrival Delay': mean_arr_array[i], 'Mean Departure Delay': mean_dep_array[i],
                        'LONGITUDE': longitude_coor[i], 'LATITUDE': latitude_coor[i]}
    i += 1
```

```
In [ ]: # Arrival Delay Figure
fig, ax = plt.subplots()
# s=0.2 specifies the size
for key in key_array:
    if airport_stats[key]['Mean Arrival Delay'] <= 0:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=0.1, c='blue')
    elif airport_stats[key]['Mean Arrival Delay'] <= 15:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=1, c='orange')
    elif airport_stats[key]['Mean Arrival Delay'] <= 30:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=2, c='green')
    else:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=5, c='red')
ax.set_title('Delay of Arrivals at USA Airport Locations')
ax.set_xlabel('Longitude')
ax.set_ylabel('Latitude');
fig.set_dpi(200)
```

```
In [ ]: # Departure Delay Figure
fig, ax = plt.subplots()
# s=0.2 specifies the size
for key in key_array:
    if airport_stats[key]['Mean Departure Delay'] <= 0:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=0.1, c='green')
    elif airport_stats[key]['Mean Departure Delay'] <= 15:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=1, c='yellow')
    elif airport_stats[key]['Mean Departure Delay'] <= 30:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=2, c='orange')
    else:
        ax.scatter(airport_stats[key]['LONGITUDE'], airport_stats[key]['LATITUDE'], s=5, c='red')
ax.set_title('Delay of Departures at USA Airport Locations')
ax.set_xlabel('Longitude')
ax.set_ylabel('Latitude');
fig.set_dpi(200)
```



```
In [ ]: # [['AIRLINE']] shows dataframe of airlines
df = flights_data[['AIRLINE', 'DEPARTURE_DELAY']]

# adding .AIRLINE at the end creates the rows to be CALLABLE
airlines = df.AIRLINE
```

```
In [ ]: # The following calculates the mean of all columns
# and groups all similar names i.e. airlines under the same group
dep_delay_mean = df.groupby(airlines).mean()
```

```
In [ ]: # .sort_values will sort alphabetically the column specified in a csv file
key_array = airlines_data.sort_values(['IATA_CODE'])
# adding .IATA_CODE at the end creates the rows to be CALLABLE
key_array = key_array.IATA_CODE

# .values allows for the mean values of departure delay to be callable
mean_val_array = dep_delay_mean[['DEPARTURE_DELAY']].values[:,0]
# .around will round up values throughout the entire array
mean_val_array = np.around(mean_val_array, decimals=2)
```

```
In [ ]: # The LHS creates a new column, restoring it with the cancelled flights from flights.c
df['CANCELLED'] = flights_data['CANCELLED']
# remove departure delay
df = df.drop(['DEPARTURE_DELAY'], axis = 1)
```

```
<ipython-input-24-dc92ab280789>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['CANCELLED'] = flights_data['CANCELLED']
```

```
In [ ]: # The following sums the total for all columns
# and groups all similar names i.e. airlines under the same group
canceled_flight_sum = df.groupby(airlines).sum()
```

```
In [ ]: # .values allows for the sum values of cancelled flights to be callable
sum_val_array = canceled_flight_sum [['CANCELLED']].values[:,0]
# .around will round up values throughout the entire array
sum_val_array = np.around(sum_val_array, decimals=2)
```

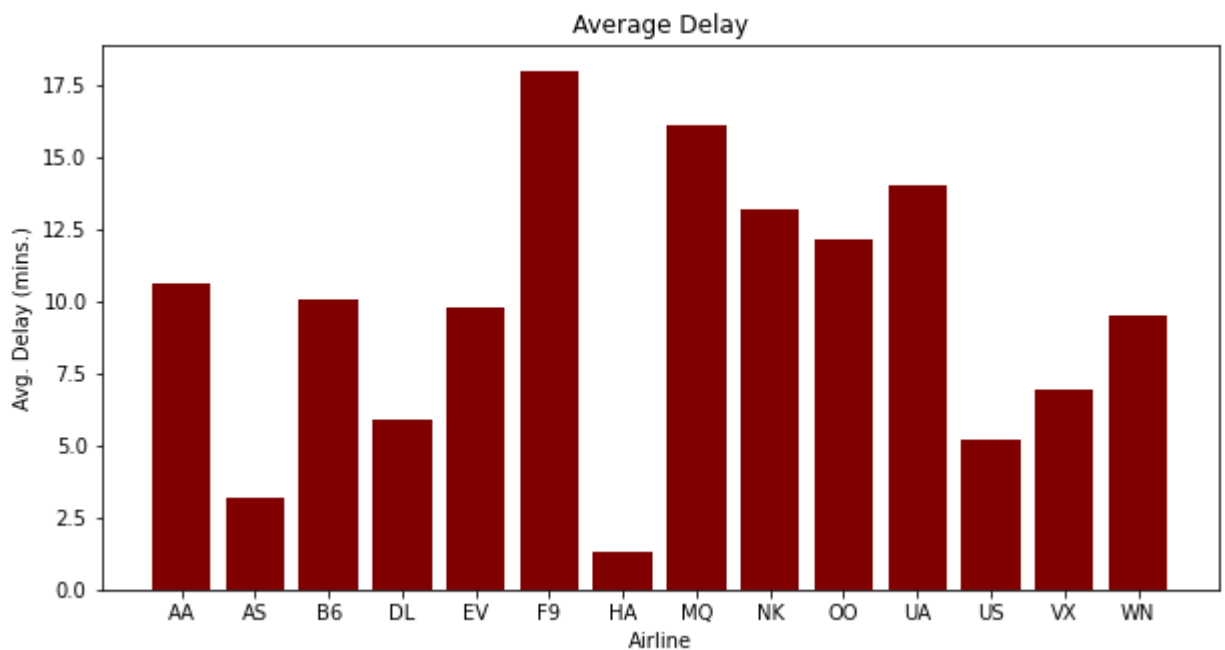
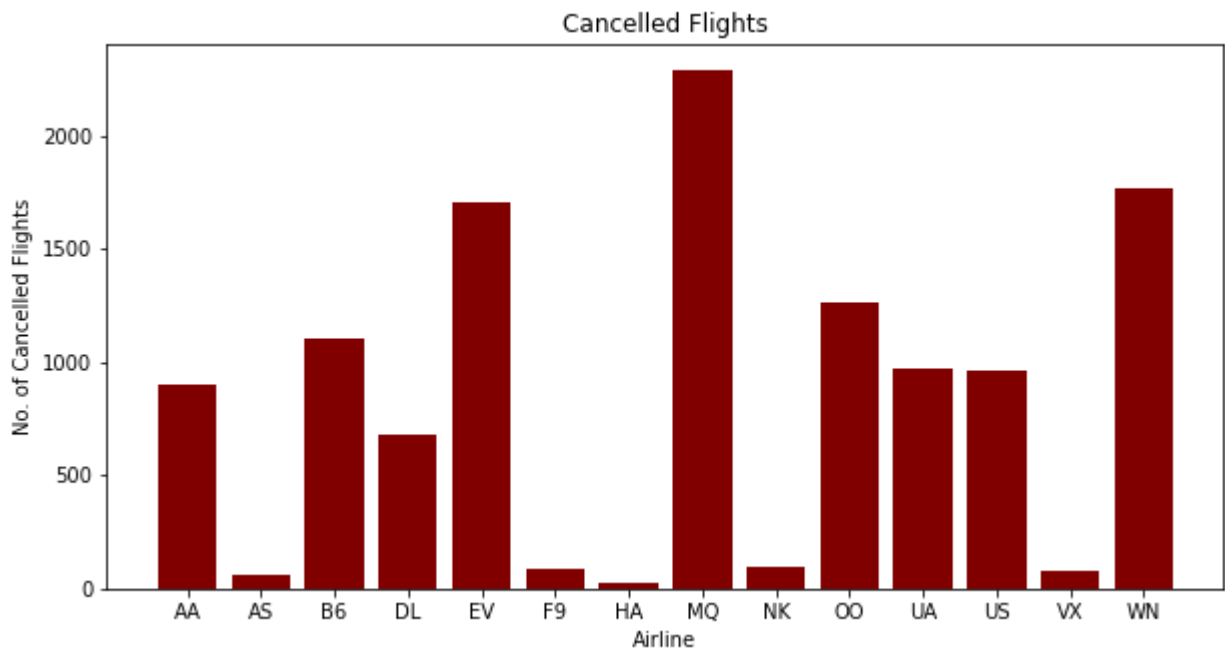
```
In [ ]: # Loop through each key and adds to an empty dictionary
# adding the appropriate stats for each key
airline_stats = {}
i = 0
for key in key_array:
    airline_stats[key] = {'Mean Departure Delay': mean_val_array[i],
                        'Cancelled Flights': sum_val_array[i]}
    i += 1
```

```
In [ ]: # Turn dictionary keys to a list
# Following can also be done list(key_arrays)
airlines_list = list(airline_stats.keys())
```

```
In [ ]: fig1 = plt.figure(figsize = (10, 5))
# creating the bar plot
plt.bar(airlines_list, sum_val_array, color = 'maroon')
plt.xlabel("Airline")
plt.ylabel("No. of Cancelled Flights")
plt.title("Cancelled Flights")
plt.show()

print()

fig2 = plt.figure(figsize = (10, 5))
# creating the bar plot
plt.bar(airlines_list, mean_val_array, color = 'maroon')
plt.xlabel("Airline")
plt.ylabel("Avg. Delay (mins.)")
plt.title("Average Delay")
plt.show()
```



```
In [ ]: # dftrain drops cols listed, keeping the cols we want to work with
dftrain = flights_data.drop(['MONTH', 'DAY', 'CANCELLATION_REASON', 'AIR_SYSTEM_DELAY',
                             'SECURITY_DELAY', 'AIRLINE_DELAY',
                             'LATE_AIRCRAFT_DELAY', 'WEATHER_DELAY'], axis = 1)

dftrain = dftrain.dropna()
# pred_input is all columns except the column being predicted
# pred is the column that will be predicted
pred_input = dftrain.drop('ARRIVAL_DELAY', axis = 1)
ydata = dftrain['ARRIVAL_DELAY']

# Train data
x_train, x_test, y_train, y_test = train_test_split(pred_input, ydata, test_size = 0.3)
```

```
In [ ]: # .info() helps us see what is not an int or float which will then be removed
# this is necessary for standard scaling of data
dftrain.info()
```

```
# Airline, origin airport, and destination airport are dropped as they are not int or
x_train = x_train.drop(['AIRLINE', 'ORIGIN_AIRPORT', 'DESTINATION_AIRPORT'], axis = 1)
x_test = x_test.drop(['AIRLINE', 'ORIGIN_AIRPORT', 'DESTINATION_AIRPORT'], axis = 1)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 457013 entries, 0 to 469967
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   AIRLINE                457013 non-null object
1   ORIGIN_AIRPORT         457013 non-null object
2   DESTINATION_AIRPORT    457013 non-null object
3   SCHEDULED_DEPARTURE    457013 non-null int64
4   DEPARTURE_TIME         457013 non-null float64
5   DEPARTURE_DELAY        457013 non-null float64
6   TAXI_OUT               457013 non-null float64
7   WHEELS_OFF             457013 non-null float64
8   SCHEDULED_TIME         457013 non-null int64
9   ELAPSED_TIME           457013 non-null float64
10  AIR_TIME               457013 non-null float64
11  DISTANCE               457013 non-null int64
12  WHEELS_ON              457013 non-null float64
13  TAXI_IN                457013 non-null float64
14  SCHEDULED_ARRIVAL      457013 non-null int64
15  ARRIVAL_TIME           457013 non-null float64
16  ARRIVAL_DELAY          457013 non-null float64
17  DIVERTED               457013 non-null int64
18  CANCELLED              457013 non-null int64
dtypes: float64(10), int64(6), object(3)
memory usage: 69.7+ MB
```

```
In [ ]: # train and test data is standardized. standard scaling is explained well here:
# https://www.digitalocean.com/community/tutorials/standardscaler-function-in-python
```


```
sc = StandardScaler()
X_train_sc = sc.fit_transform(x_train)
X_test_sc = sc.transform(x_test)
```

```
In [ ]: # models used for training
LinR = LinearRegression()
Rid = Ridge()
```

```
In [ ]: # Output of model prediction vs truth as well as MSE, MAE, Root MSE, and R2
for model, name in zip([LinR, Rid], ['Linear Regression', 'Ridge']):
    model1 = model.fit(X_train_sc, y_train)
    Y_predict = model1.predict(X_test_sc)
    print(name)
    print()
    plt.scatter(y_test, Y_predict)
    plt.title("Model Analysis")
    plt.xlabel("Truth")
    plt.ylabel("Prediction")
    plt.show()
    print('Mean Absolute Error:', mean_absolute_error(y_test, Y_predict))
    print('Mean Squared Error:', mean_squared_error(y_test, Y_predict))
    print('Root Mean Squared Error:', np.sqrt(mean_squared_error(y_test, Y_predict)))
```


```
print('R2 : ',r2_score(y_test, Y_predict))  
print()
```

Linear Regression



Mean Absolute Error: 8.068423394684439e-14
Mean Squared Error: 1.2526479352295827e-26
Root Mean Squared Error: 1.1192175549148533e-13
R2 : 1.0

Ridge



Mean Absolute Error: 0.0019259705131082696
Mean Squared Error: 6.526618509607398e-06
Root Mean Squared Error: 0.0025547247424345734
R2 : 0.9999999958602483